

# 智能体（Agent）设计与实现详解

CS181-Project

June 2, 2025

# 目录

- 1 智能体体系结构总览
- 2 评估函数详解
- 3 RandomPlayer & GreedyPlayer
- 4 MinimaxPlayer
- 5 MCTSPlayer
- 6 MinimaxPlayer
- 7 ApproximateQLearningPlayer
- 8 ApproximateQLearningPlayer
- 9 ApproximateQLearningPlayer
- 10 网络结构
- 11 训练细节
- 12 Q值与更新公式
- 13 实现细节
- 14 Network Architecture
- 15 Training Details
- 16 Q-value and Update Formula
- 17 Implementation Details

- **Player**: 抽象基类，定义所有玩家的接口
- **HumanPlayer**: 人类玩家，动作由外部输入
- **RandomPlayer**: 随机选择动作，作为基线
- **GreedyPlayer**: 贪心选择当前最优动作
- **MinimaxPlayer**: 极大极小搜索，考虑对手反应
- **MCTSPlayer**: 蒙特卡洛树搜索，平衡探索与利用
- **ApproximateQLearningPlayer**: 特征线性Q学习
- **Neural\_ApproximateQLearningPlayer**: 深度Q学习（神经网络）

用于**Greedy/Minimax**智能体，衡量棋盘状态优劣。

- 计算每个己方棋子到目标区所有空位的最大欧氏距离
- 所有棋子都到达目标区时有惩罚（-20）
- 返回值乘以-1，使得距离越小分数越高

$$\text{val} = - \sum_{p \in P_{\text{self}}} \begin{cases} \max_{g \in G_{\text{empty}}} \sqrt{(x_p - x_g)^2 + (y_p - y_g)^2}, & \text{if } G_{\text{empty}} \neq \emptyset \\ -20, & \text{if } G_{\text{empty}} = \emptyset \end{cases}$$

$$\text{dist}(p, g) = \sqrt{(x_p - x_g)^2 + (y_p - y_g)^2}$$

# 评估函数 evaluation\_MCTS

用于**MCTS**和强化学习智能体，综合多项指标。

- 目标区棋子数量奖励（40%）
- 距离评分（30%）
- 阶段性奖励（递进式）
- 家中棋子惩罚

$$\text{goal\_score} = 0.4 \times 1000 \times \frac{\text{pieces\_in\_goal}}{4}$$

$$\text{distance\_score} = 0.3 \times 300 \times \left( 1 - \frac{\text{normalized\_dist\_sum}}{N} \right)$$

stage\_bonus = 分阶段奖励

$$\text{home\_penalty} = -100 \times \frac{\text{pieces\_at\_home}}{N}$$

$$\text{final\_score} = \text{goal\_score} + \text{distance\_score} + \text{stage\_bonus} + \text{home\_penalty}$$

## RandomPlayer

- 完全随机选择动作
- 用于基线对比

## GreedyPlayer

- 遍历所有动作，模拟后用evaluation评估
- 选择分数最高的动作
- 只考虑当前一步

# MinimaxPlayer (极大极小搜索)

- 递归搜索到指定深度或超时
- 己方回合最大化分数，对手回合最小化分数
- 使用Alpha-Beta剪枝优化搜索效率
- 叶节点用evaluation评估
- 支持**Local Search** (局部搜索) 技术，进一步提升效率

$$\text{Minimax}(s, d) = \begin{cases} \max_{a \in A(s)} \text{Minimax}(s', d - 1), & \text{if maximizer} \\ \min_{a \in A(s)} \text{Minimax}(s', d - 1), & \text{if minimizer} \end{cases}$$

**Local Search (局部搜索)** 是一种在博弈树搜索中减少分支、提升效率的启发式方法。

- 标准Minimax每层枚举所有己方棋子的所有动作，分支数极大。
- Local Search只为每个己方棋子挑选最优（或前k优）动作，大幅减少分支。
- 实现方式：对每个己方棋子，模拟所有可行动作，用evaluation函数评估，仅保留分数最高的动作。
- 这样能让搜索更深，提升决策速度，同时保证整体推进。

伪代码：

$$A_{\text{local}}(s) = \bigcup_{p \in P_{\text{self}}} \left\{ \arg \max_{a \in A_p(s)} \text{evaluation}(s, a) \right\}$$

其中 $A_p(s)$ 为棋子 $p$ 所有可行动作。



# Local Search的优缺点

## 优点:

- 极大减少分支数，提升搜索深度和效率
- 保证每个己方棋子都被推进，避免只关注部分棋子
- 实现简单，易于集成到Minimax框架

## 缺点:

- 可能丢失全局最优解（只考虑局部最优）
- 对evaluation函数的准确性依赖较强

## 四阶段流程

- ① **Selection:** 用UCB选择最优子节点
- ② **Expansion:** 扩展新子节点, 优先朝目标方向
- ③ **Simulation:** 90%概率用启发式, 10%随机, 最多30步
- ④ **Backpropagation:** 回传模拟结果, 更新访问次数和累计价值

**UCB (Upper Confidence Bound) 公式:**

$$UCB = \frac{Q}{N} + c\sqrt{\frac{\ln N_{parent}}{N}} + \text{strategy\_score}$$

- $Q$ : 累计价值
- $N$ : 访问次数
- $c$ : 探索系数 (前期大, 后期小)
- $\text{strategy\_score}$ : 方向性、跳跃奖励等

# MCTS动作优先级与模拟

- 优先减少到目标距离的动作
- 奖励跳跃，惩罚后退
- 模拟阶段90%概率选择朝目标方向动作
- 最终动作选择综合胜率、访问率、方向性

$$\text{score} = 0.4 \times \text{win\_ratio} + 0.2 \times \text{visit\_ratio} + 0.4 \times \text{direction\_score}$$

# MinimaxPlayer (Minimax Search)

- Recursively searches to a specified depth or until timeout.
- Maximizes the score on the player's turn, minimizes on the opponent's turn.
- Uses Alpha-Beta pruning to improve search efficiency.
- Uses the evaluation function at leaf nodes.
- Supports **Local Search** to further improve efficiency.

$$\text{Minimax}(s, d) = \begin{cases} \max_{a \in A(s)} \text{Minimax}(s', d - 1), & \text{if maximizer} \\ \min_{a \in A(s)} \text{Minimax}(s', d - 1), & \text{if minimizer} \end{cases}$$

# Local Search in Minimax

**Local Search** is a heuristic method to reduce the branching factor and improve efficiency in game tree search.

- Standard Minimax enumerates all possible moves for all pawns at each layer, resulting in a huge branching factor.
- Local Search selects only the best (or top-k) move(s) for each pawn, greatly reducing the branching factor.
- Implementation: For each pawn, simulate all possible moves, evaluate them using the evaluation function, and keep only the move with the highest score.
- This allows deeper search and faster decision-making, while ensuring all pawns are advanced.

**Pseudocode:**

$$A_{\text{local}}(s) = \bigcup_{p \in P_{\text{self}}} \left\{ \arg \max_{a \in A_p(s)} \text{evaluation}(s, a) \right\}$$

where  $A_p(s)$  is the set of all possible moves for pawn  $p$ .

# Advantages and Disadvantages of Local Search

## Advantages:

- Greatly reduces the branching factor, allowing deeper and more efficient search.
- Ensures all pawns are advanced, avoiding focus on only a few pawns.
- Simple to implement and easy to integrate into the Minimax framework.

## Disadvantages:

- May miss the global optimum (only considers local optima).
- Relies heavily on the accuracy of the evaluation function.

## 特征设计

- 状态特征：目标区棋子比例、平均距离
- 动作特征：距离改善、方向性、跳跃、到达目标、后退、离家

$$Q(s, a) = \sum_i w_i \cdot f_i(s, a)$$

## 权重示例：

- pieces\_in\_goal: 2000
- avg\_distance: -800
- distance\_improvement: 500
- direction: 300
- is\_jump: 100
- reaches\_goal: 1500
- is\_backwards: -1000
- leaves\_home: 200



**TD(0)更新公式:**

$$w_i \leftarrow w_i + \alpha \cdot (r + \gamma \max_{a'} Q(s', a') - Q(s, a)) \cdot f_i(s, a)$$

**奖励设计**

- 胜利奖励:  $3000 + 500 \times \text{pieces\_in\_goal}$
- 目标进展奖励:  $300 \times 2^{\text{当前目标棋子数}}$
- 距离改善奖励: 前期100, 后期200
- 跳跃奖励: 前期200, 中期50, 后期0
- 后退惩罚: 前期-200, 后期-500

## Feature-based Q-Learning:

- Uses a linear combination of hand-crafted features to estimate  $Q(s, a)$ .
- Features are divided into state features and action features.

$$Q(s, a) = \sum_i w_i \cdot f_i(s, a)$$

# Neural\_ApproximateQLearningPlayer 总结

**核心思想:** 使用深度神经网络 (DNN) 来近似Q值函数  $Q(s, a)$ , 结合经验回放和目标网络进行训练。

## 网络结构 (NeuralFeatureExtractor):

- 输入:

- 棋盘状态 (Board State):  $4 \times N \times N$  张量 (己方棋子, 对方棋子, 目标区, 初始区)
- 动作 (Action): 4维向量 (start\_x, start\_y, end\_x, end\_y)
- 棋盘编码: 3层卷积网络 (Conv2d)  $\rightarrow$  Flatten
- 动作编码: 1层全连接网络 (Linear)
- 特征融合: 拼接棋盘和动作特征  $\rightarrow$  3层全连接网络 (Linear)
- 输出: 单个Q值  $Q(s, a)$

## Q值函数:

$$Q(s, a; \theta) = \text{DNN}(\text{encode\_board}(s), \text{encode\_action}(a); \theta)$$

其中  $\theta$  是神经网络的参数。

## 训练与更新 (DQN思想):

# State Features (Examples)

## State features:

$$\text{pieces\_in\_goal} = \frac{\text{Number of player's pieces in goal area}}{4}$$

$$\text{avg\_distance} = \frac{1}{4D_{\max}} \sum_{p \notin G} |x_p - x_c| + |y_p - y_c|$$

- $G$ : goal area,  $(x_c, y_c)$ : goal center,  $D_{\max}$ : max possible distance

# Action Features (Examples)

## Action features:

$$\text{distance\_improvement} = \frac{d_{\text{start}} - d_{\text{end}}}{D_{\text{max}}}$$

$$\text{direction} = \begin{cases} \frac{(x_{\text{end}} - x_{\text{start}}) + (y_{\text{end}} - y_{\text{start}})}{2B}, & \text{if RED} \\ \frac{(x_{\text{start}} - x_{\text{end}}) + (y_{\text{start}} - y_{\text{end}})}{2B}, & \text{otherwise} \end{cases}$$

$$\text{is\_jump} = \begin{cases} 1, & \text{if jump move} \\ 0, & \text{otherwise} \end{cases}$$

$$\text{reaches\_goal} = \begin{cases} 1, & (x_{\text{end}}, y_{\text{end}}) \in G \\ 0, & \text{otherwise} \end{cases}$$

## Additional action features:

$$\text{is\_backwards} = \begin{cases} 1, & \text{if move is backwards} \\ 0, & \text{otherwise} \end{cases}$$

$$\text{leaves\_home} = \begin{cases} 1, & (x_{\text{start}}, y_{\text{start}}) \in H \text{ and } (x_{\text{end}}, y_{\text{end}}) \notin H \\ 0, & \text{otherwise} \end{cases}$$

- $H$ : home area

# Example Feature Weights

## Example weights:


- pieces\_in\_goal: 2000
- avg\_distance: -800
- distance\_improvement: 500
- direction: 300
- is\_jump: 100
- reaches\_goal: 1500
- is\_backwards: -1000
- leaves\_home: 200

$$Q(s, a) = \sum_{i=1}^n w_i \cdot f_i(s, a)$$

$$\begin{aligned} Q(s, a) = & w_1 \cdot \text{pieces\_in\_goal} \\ & + w_2 \cdot \text{avg\_distance} \\ & + w_3 \cdot \text{distance\_improvement} \\ & + w_4 \cdot \text{direction} \\ & + w_5 \cdot \text{is\_jump} \\ & + w_6 \cdot \text{reaches\_goal} \\ & + w_7 \cdot \text{is\_backwards} \\ & + w_8 \cdot \text{leaves\_home} \end{aligned}$$



# 信息流结构图



- 经验回放 (Replay Buffer) : deque, 容量10000
- 每步采样32个样本进行小批量训练
- 目标网络每1000步同步一次
- 损失函数: 均方误差 (MSE)
- 优化器: Adam
- $\epsilon$ -greedy策略动态调整探索率
- 过滤掉反向动作, 避免无效循环
- 终局阶段 (仅剩2个未进目标区棋子) 采用Minimax逻辑

# Q值与更新公式

$$Q(s, a) = \text{network}(s, a)$$

$$\text{Loss} = \frac{1}{N} \sum_{i=1}^N \left( Q(s_i, a_i) - \left[ r_i + \gamma \max_{a'} Q_{\text{target}}(s'_i, a') \right] \right)^2$$

- $Q_{\text{target}}$ : 目标网络
- $r_i$ : 即时奖励,  $\gamma$ : 折扣因子

- 支持GPU加速（自动检测CUDA）
- 支持模型保存与加载
- 训练时动态调整 $\epsilon$ ，后期更倾向利用
- 经验回放提升样本利用率，目标网络提升训练稳定性
- 支持自我对弈训练和与Minimax等对手对弈


# 智能体（Agent）设计与实现详解

CS181-Project

June 2, 2025

- **Input:** Board state (4 channels,  $N \times N$ ) and action (4-dimensional vector)
- **Board encoder:** 3 convolutional layers (Conv2d)
- **Action encoder:** 1 fully connected (FC) layer
- **Fusion:** 3 FC layers with ReLU and Dropout
- **Output:** Q-value

# Information Flow Diagram



# Training Details

- Experience replay buffer (deque, size 10000)
- Mini-batch update (batch size 32)
- Target network updated every 1000 steps
- Loss: Mean Squared Error (MSE)
- Optimizer: Adam
- $\epsilon$ -greedy policy with dynamic exploration rate
- Filters out reverse actions to avoid oscillation
- Uses Minimax logic for endgame (when only 2 pieces not in goal)



# Q-value and Update Formula

$$Q(s, a) = \text{network}(s, a)$$

$$\text{Loss} = \frac{1}{N} \sum_{i=1}^N \left( Q(s_i, a_i) - \left[ r_i + \gamma \max_{a'} Q_{\text{target}}(s'_i, a') \right] \right)^2$$

- $Q_{\text{target}}$ : target network
- $r_i$ : immediate reward,  $\gamma$ : discount factor

# Implementation Details

- Supports GPU acceleration (auto-detects CUDA)
- Supports model saving and loading
- Dynamically adjusts  $\epsilon$  during training (more exploitation in late stage)
- Experience replay improves sample efficiency; target network stabilizes training
- Supports self-play and training against Minimax or other agents

# 智能体对比总结

类型	策略核心	优点	缺点
Random	随机	实现极简	无智能
Greedy	贪心	推进快，简单	只看一步
Minimax	博弈树	考虑对手	分支爆炸
MCTS	蒙特卡洛树	探索与利用平衡	计算量大
AQL	特征Q学习	可解释性强	特征有限
NAQL	深度Q学习	泛化强	训练慢